




Siddhant Mishra-Sharma (MIT/AI FI) Summer School



170

30

Controlling compression and disentanglement:  $\beta$ -VAEs



$$-\text{ELBO} = -\left\langle \log p_{\vartheta}(x \mid z) \right\rangle_{q_{\varphi}} + \textcolor{red}{\beta} \cdot D_{\text{KL}}\left(q_{\varphi}(z \mid x) \parallel p(z)\right)$$





Rate



Disortion

If the data-generating process is associated with a principled noise model, by using it (the *likelihood*) as the reconstruction loss we are aiming to reconstruct the mean data.

$$\log p(x \mid z; x') = -\frac{1}{2} \left( \frac{x - x'}{\sigma} \right)^2 + \log \left( \frac{1}{\sigma \sqrt{2\pi}} \right)$$

- Larger  $\sigma$ : More of the data variation is attributed to the likelihood  $\rightarrow$  larger " $\beta$ ", more compression
- Smaller  $\sigma$ : Latents  $z$  try to capture more of the variation in the data (e.g. small perceptual features)


# Controlling compression and disentanglement: $\beta$ -VAEs

$$-\text{ELBO} = \underbrace{-\left\langle \log p_{\vartheta}(x \mid z) \right\rangle_{q_{\varphi}}}_{\text{Distortion}} + \underbrace{\beta \cdot D_{\text{KL}} \left( q_{\varphi}(z \mid x) \parallel p(z) \right)}_{\text{Rate}}$$

If the data-generating process is associated with a principled noise model, by using it (the *likelihood*) as the reconstruction loss we are aiming to reconstruct the mean data.


$$\log p(x \mid z; x') = -\frac{1}{2} \left( \frac{x - x'}{\sigma} \right)^2 + \log \left( \frac{1}{\sigma \sqrt{2\pi}} \right)$$

- Larger  $\sigma$ : More of the data variation is attributed to the likelihood  $\rightarrow$  larger " $\beta$ ", more compression
- Smaller  $\sigma$ : Latents  $z$  try to capture more of the variation in the data (e.g. small perceptual features)



# Tutorials 1 and 2: variational inference and VAEs

Lead: Carol Cuesta-Lazaro



- Implement the ELBO objective for **variational inference**
- Construct a VAE and use it to build a **generative model of galaxy images** using samples from the HST COSMOS dataset
- Boilerplate code for training/reconstruction/sampling for quick iteration
- Experiment with trade-offs between pure **reconstruction** and a latent space **regularization**

[Mandelbaum et al; <https://zenodo.org/record/3242143>]